

# S Cloudfront & AWS Global Accelerator

## Amazon CloudFront

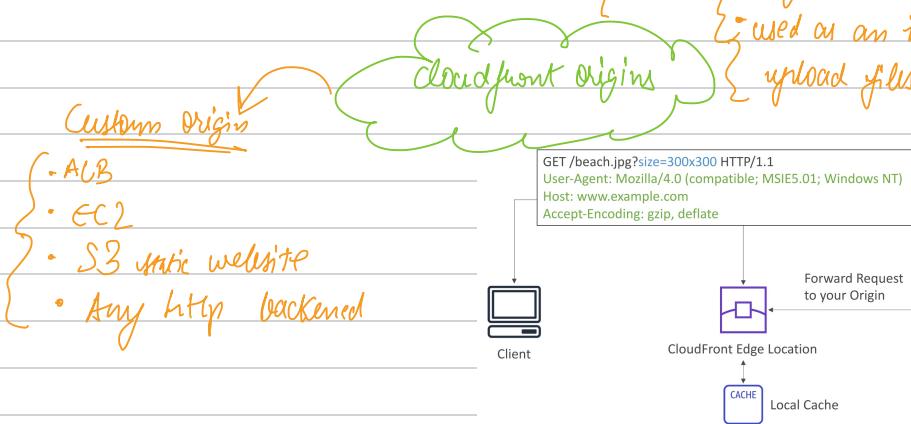
- Content Delivery Network (CDN)
- Improves read performance, content is cached at the edge
- Improves users experience
- 216 Point of Presence globally (edge locations)
- DDoS protection (because worldwide), integration with Shield, AWS Web Application Firewall

→ *edge distributed*



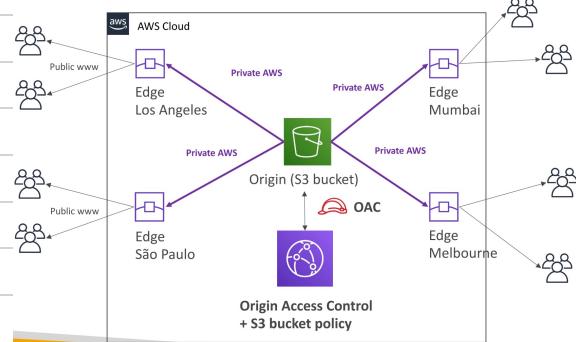
## S3 Bucket

- { for distributing & caching
- { origin access control (OAC)
- { used as an ingress to upload files to S3



## CloudFront vs S3 Cross Region Replication

- CloudFront:
  - Global Edge network
  - Files are cached for a TTL (maybe a day)
  - Great for static content that must be available everywhere
- S3 Cross Region Replication:
  - Must be setup for each region you want replication to happen
  - Files are updated in near real-time
  - Read only
  - Great for dynamic content that needs to be available at low-latency in few regions



→ We can use CloudFront step without making the bucket public, by using origin access control

go to cloudfront > distributions > create distribution

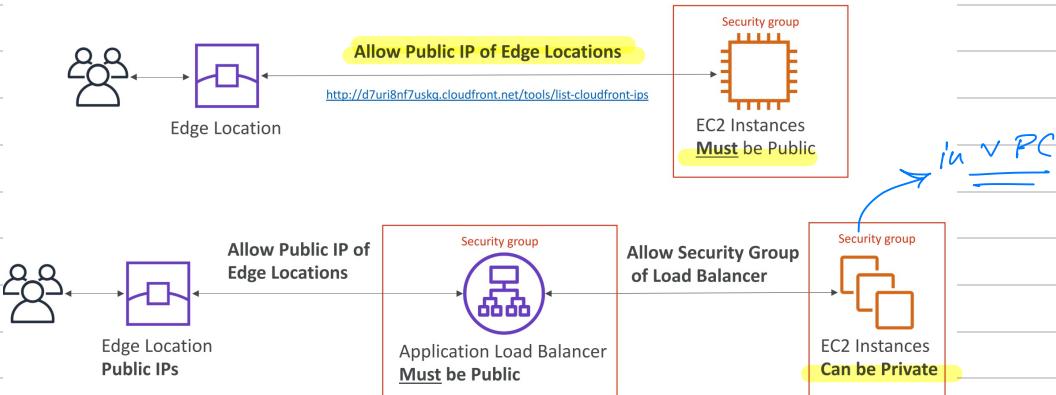
new < Origins control < Origin access control < name  
to update < selecting  
bucket policy.

✓  
dynamic root  
object = index.html .

→ adding bucket policy for  
CloudFront S3 S3  
Edit origins just go to  
edit distribution & copy  
policy! paste S3 to S3  
bucket!

Once deployed, just  
open the domain in cloudfront.

## CloudFront – ALB or EC2 as an origin



# CloudFront Geo Restriction

- You can restrict who can access your distribution
  - **Allowlist:** Allow your users to access your content only if they're in one of the countries on a list of approved countries.
  - **Blocklist:** Prevent your users from accessing your content if they're in one of the countries on a list of banned countries.
- The "country" is determined using a 3<sup>rd</sup> party Geo-IP database
- Use case: Copyright Laws to control access to content

distribution >  
security >  
cloudfront geo  
restrictions.

→ pricing is cloudfront is different all around the world.

more the

data out of cloudfront

lower the price

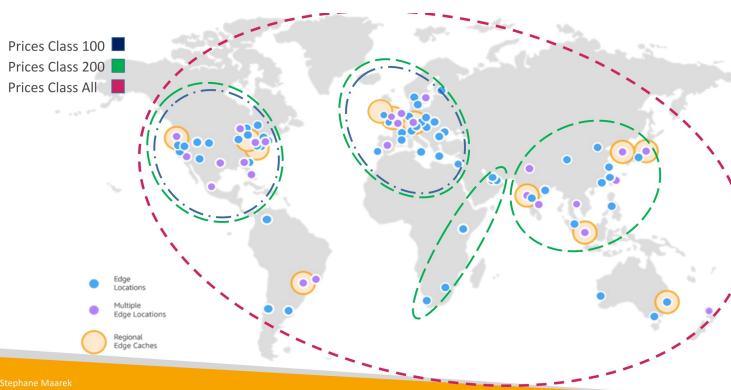
min → US Mexico & Canada

med → India

## CloudFront – Price Classes

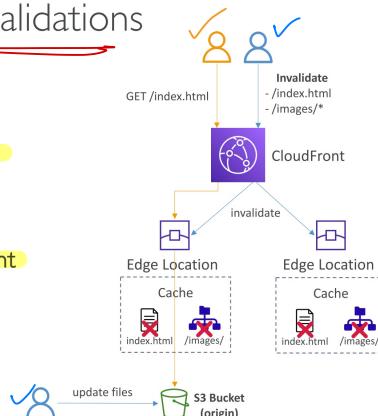
- You can reduce the number of edge locations for cost reduction
- Three price classes:
  1. **Price Class All:** all regions – best performance
  2. **Price Class 200:** most regions, but excludes the most expensive regions
  3. **Price Class 100:** only the least expensive regions

Edge Locations Included Within	United States, Mexico, & Canada	Europe & Israel	South Africa, Kenya, & Middle East	South America	Japan	Australia & New Zealand	Hong Kong, Philippines, Singapore, South Korea, Taiwan, & Thailand	India
Price Class All	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Price Class 200	Yes	Yes	Yes	x	Yes	x	Yes	Yes
Price Class 100	Yes	Yes	x	x	x	x	x	x



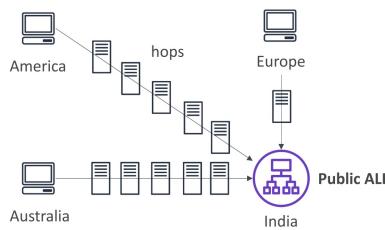
## CloudFront – Cache Invalidation

- In case you update the back-end origin, CloudFront doesn't know about it and will only get the refreshed content after the TTL has expired
- However, you can force an entire or partial cache refresh (thus bypassing the TTL) by performing a **CloudFront Invalidation**
- You can invalidate all files (\*) or a special path (/images/\*)



## Global users for our application

- You have deployed an application and have global users who want to access it directly.
- They go over the public internet, which can add a lot of latency due to many hops
- We wish to go as fast as possible through AWS network to minimize latency



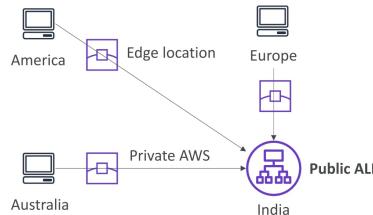
• Unicast IP: one server holds one IP address  
• Anycast IP: all servers hold the same IP address and the client is routed to the nearest one

*this concept is used by global accelerator.*



## AWS Global Accelerator

- Leverage the AWS internal network to route to your application
- 2 Anycast IP are created for your application
- The Anycast IP send traffic directly to Edge Locations
- The Edge locations send the traffic to your application



→ works with EC2, Elastic IPs, ALB, NLB, public or private

- Consistent Performance
  - Intelligent routing to lowest latency and fast regional failover
  - No issue with client cache (because the IP doesn't change)
  - Internal AWS network
- Health Checks
  - Global Accelerator performs a health check of your applications
  - Helps make your application global (failover less than 1 minute for unhealthy)
  - Great for disaster recovery (thanks to the health checks)
- Security
  - only 2 external IP need to be whitelisted
  - DDoS protection thanks to AWS Shield

→ global service  
but must specify  
US West (Oregon)

→ if endpoint is ALB then  
will also leverage its  
health checks.

## AWS Global Accelerator vs CloudFront

- They both use the AWS global network and its edge locations around the world
- Both services integrate with AWS Shield for DDoS protection.
- CloudFront
  - Improves performance for both cacheable content (such as images and videos)
  - Dynamic content (such as API acceleration and dynamic site delivery)
  - Content is served at the edge
- Global Accelerator
  - Improves performance for a wide range of applications over TCP or UDP
  - Proxying packets at the edge to applications running in one or more AWS Regions.
  - Good fit for non-HTTP use cases, such as gaming (UDP), IoT (MQTT), or Voice over IP
  - Good for HTTP use cases that require static IP addresses
  - Good for HTTP use cases that required deterministic, fast regional failover

content is cached  
at edge

the packets make  
it do the  
application but  
by using  
internal network

create a application on EC2 instance which exposes as  
HTTP endpoint. in 2 regions.  
Security group (DJ1137) for HTTP

create global accelerator.

Add listeners { port  
protocol  
client affinity { like sticky sessions}

Add endpoint groups { add constraints for the EC2 instances