

## § Containers on AWS

Q What is a docker?

→ It's a software development platform to deploy apps. Apps are packaged into containers that can run on any OS.

→ good for microservices-architecture & life-and-shift apps from On-premises to AWS cloud.  
app on

→ can run a docker container on any OS and more than one instance.

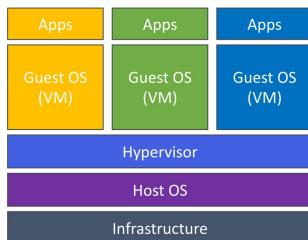
Docker images can be stored in

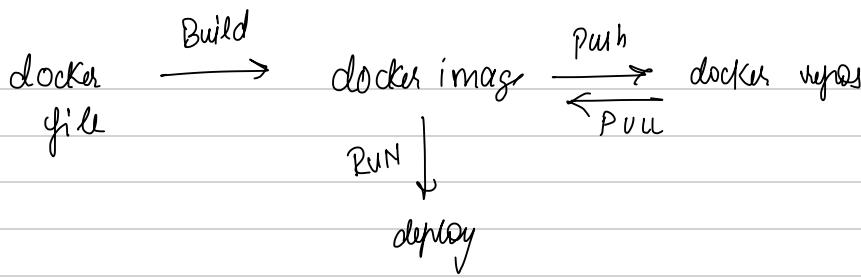
Docker hub  
{ Public repos }  
can find reuse images of  
many Technologies & OS

Amazon ECR  
{ Public (ECR public gallery) }  
Private

## Docker vs. Virtual Machines

- Docker is "sort of" a virtualization technology, but not exactly
- Resources are shared with the host => many containers on one server





## Docker Containers Management on AWS

- ① Amazon Elastic Container Service (Amazon ECS)
  - Amazon's own container platform
- ② Amazon Elastic Kubernetes Service (Amazon EKS)
  - Amazon's managed Kubernetes (open source)
- ③ AWS Fargate
  - Amazon's own Serverless container platform
  - Works with ECS and with EKS
- ④ Amazon ECR:
  - Store container images



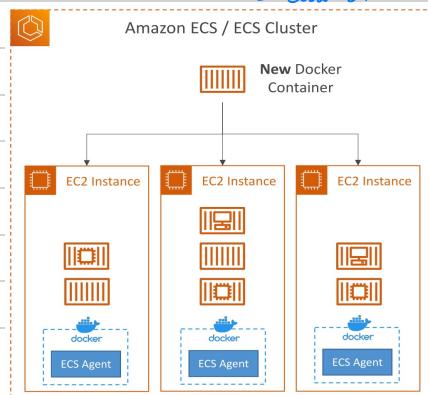
### ① Amazon ECS

{ Elastic Container Service }

- Launch docker on AWS = Launch ECS tasks on ECS cluster.

- (a) ECS launches type → provision & maintain EC2 instances.
- Each EC2 instance must run "ECS Agent" to register into ECS cluster

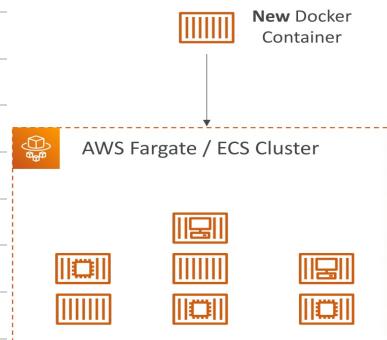
- Once done AWS takes care of starting, stopping & distributing the containers.



(3) Fargate launch type → no need to provision EC2 instances  
 ↳ underlying infrastructure is on EC2 but no need to provision ourselves

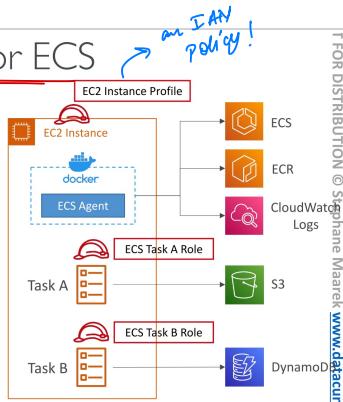
### T.C. Userless!

- Create tasks & AWS just runs them based on the CPU / RAM needed
- it will scale automatically.



## Amazon ECS – IAM Roles for ECS

- **EC2 Instance Profile (EC2 Launch Type only):**
  - Used by the ECS agent
  - Makes API calls to ECS service
  - Send container logs to CloudWatch Logs
  - Pull Docker image from ECR
  - Reference sensitive data in Secrets Manager or SSM Parameter Store
- **ECS Task Role:**
  - Allows each task to have a specific role
  - Use different roles for the different ECS Services you run
  - Task Role is defined in the task definition

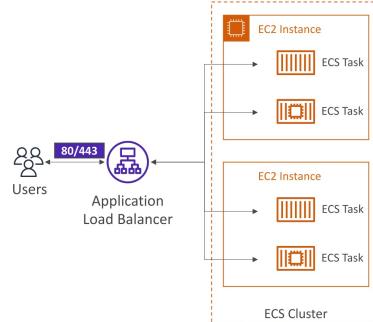


FOR DISTRIBUTION © Sebastian Maarek [www.datacur](http://www.datacur)

- ① EC2 instance profile role {instance level}
- ② ECS task role. {task level}

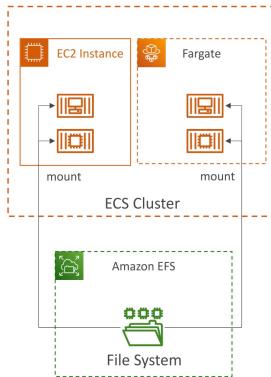
## Amazon ECS – Load Balancer Integrations

- Application Load Balancer supported and works for most use cases
- Network Load Balancer recommended only for high throughput / high performance use cases, or to pair it with AWS Private Link
- Classic Load Balancer supported but not recommended (no advanced features – no Fargate)



# Amazon ECS – Data Volumes (EFS)

- Mount EFS file systems onto ECS tasks
- Works for both EC2 and Fargate launch types
- Tasks running in any AZ will share the same data in the EFS file system
- Fargate + EFS = Serverless { ultimate }
- Use cases: persistent multi-AZ shared storage for your containers
- Note:
  - Amazon S3 cannot be mounted as a file system



① first create a ECS cluster, set vpc w/ a gateway & this will create a Fargate cluster { if EC2 launch type then can Auto Scaling Groups is also to be created. }

② To create a task go to task definition  
↳ infrastructure requirements

task role is for container to interact with other services of AWS.

## ECS Service Auto Scaling

- Automatically increase/decrease the desired number of ECS tasks
- Amazon ECS Auto Scaling uses AWS Application Auto Scaling
  - ECS Service Average CPU Utilization
  - ECS Service Average Memory Utilization - Scale on RAM
  - ALB Request Count Per Target – metric coming from the ALB
- Target Tracking – scale based on target value for a specific CloudWatch metric
- Step Scaling – scale based on a specified CloudWatch Alarm
- Scheduled Scaling – scale based on a specified date/time (predictable changes)
- ECS Service Auto Scaling (task level) ≠ EC2 Auto Scaling (EC2 instance level)
- Fargate Auto Scaling is much easier to setup (because Serverless)



↳ metrics

↳ scaling methods

# EC2 Launch Type – Auto Scaling EC2 Instances

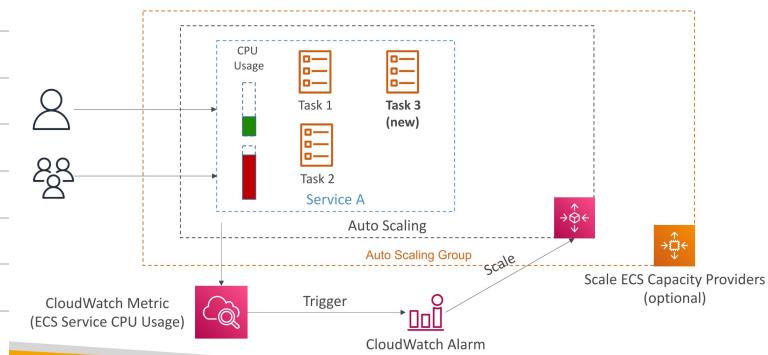
- Accommodate ECS Service Scaling by adding underlying EC2 Instances

- Auto Scaling Group Scaling

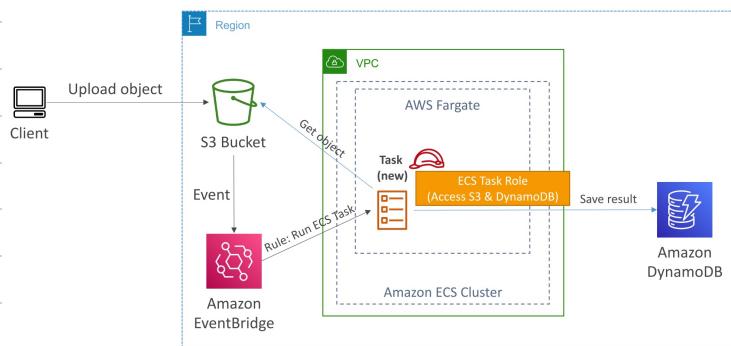
- Scale your ASG based on CPU Utilization
- Add EC2 instances over time

- ECS Cluster Capacity Provider

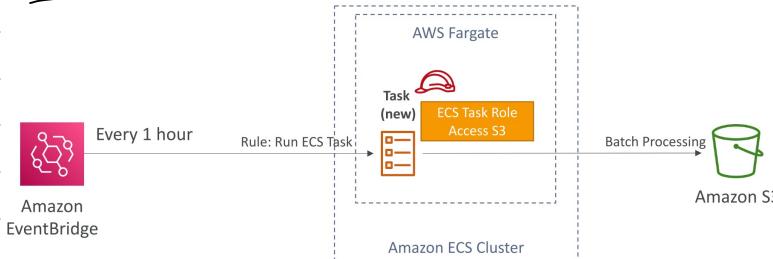
- Used to automatically provision and scale the infrastructure for your ECS Tasks
- Capacity Provider paired with an Auto Scaling Group
- Add EC2 Instances when you're missing capacity (CPU, RAM...)



\* serverless solution to process data & add them to a database.  
ECS tasks invoked by Event Bridge

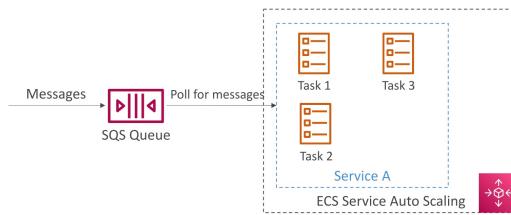


## \* Example 2



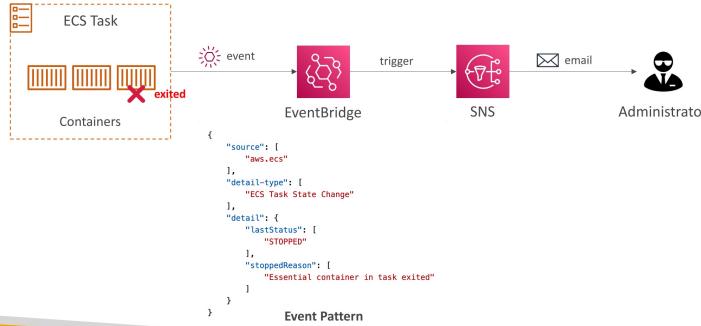
## \* Example 3

### ECS – SQS Queue Example



## \* Example 3

### ECS – Intercept Stopped Tasks using EventBridge

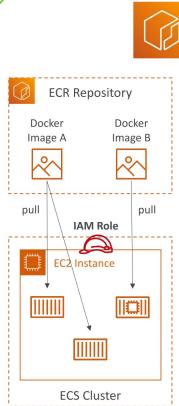


hence Event Bridge allows us to understand the lifecycle of ECS containers

## Amazon ECR

→ storing docker images ✓

- ECR = Elastic Container Registry
- Store and manage Docker images on AWS
- Private and Public repository (Amazon ECR Public Gallery <https://gallery.ecr.aws>)
- Fully integrated with ECS, backed by Amazon S3
- Access is controlled through IAM (permission errors => policy)
- Supports image vulnerability scanning, versioning, image tags, image lifecycle, ...

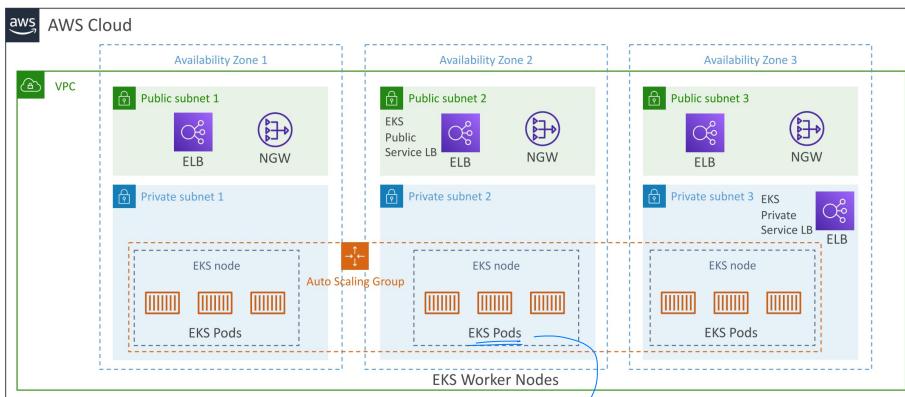


My permission denied  
is due to IAM

## Amazon EKS Overview



- Amazon EKS = Amazon Elastic Kubernetes Service
- It is a way to launch managed Kubernetes clusters on AWS
- Kubernetes is an open-source system for automatic deployment, scaling and management of containerized (usually Docker) application
- It's an alternative to ECS, similar goal but different API
- EKS supports EC2 if you want to deploy worker nodes or Fargate to deploy serverless containers
- Use case: if your company is already using Kubernetes on-premises or in another cloud, and wants to migrate to AWS using Kubernetes
- Kubernetes is cloud-agnostic (can be used in any cloud – Azure, GCP...)
- For multiple regions, deploy one EKS cluster per region
- Collect logs and metrics using CloudWatch Container Insights



→ similar to tasks

# Amazon EKS – Node Types

- Managed Node Groups *AWS managed*
  - Creates and manages Nodes (EC2 instances) **for you**
  - Nodes are part of an ASG managed by EKS
  - Supports **On-Demand or Spot Instances**
- Self-Managed Nodes
  - **Nodes created by you** and registered to the EKS cluster and managed by an ASG
  - You can **use prebuilt AMI** - Amazon EKS Optimized AMI
  - Supports **On-Demand or Spot Instances**
- AWS Fargate
  - **No maintenance** required; no nodes managed

# Amazon EKS – Data Volumes

- Need to specify **StorageClass** manifest on your EKS cluster
- Leverages a **Container Storage Interface (CSI)** compliant driver

\* *Fargate only works with EFS*

- Support for...
- Amazon EBS
- Amazon EFS (works with Fargate)
- Amazon FSx for Lustre
- Amazon FSx for NetApp ONTAP



# AWS App Runner

{ Kinda bootstrapping  
the deployment process }

- Fully managed service that makes it easy to deploy web applications and APIs at scale
- No infrastructure experience required
- Start with your source code or container image ①
- Automatically builds and deploys the web app ②
- Automatic scaling, highly available, load balancer, encryption
- VPC access support
- Connect to database, cache, and message queue services **very quickly**
- Use cases: web apps, APIs, microservices, rapid production deployments

