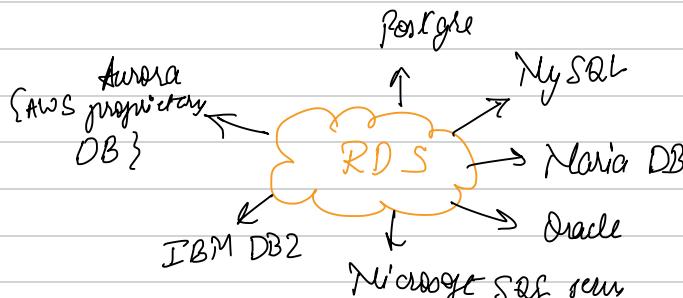


G Database

RDS { Relational Database Service }

→ AWS manage DB service for DB use as a query language

→ no need to set & provision servers to run the database, all underlying software & patch installations done by AWS itself.

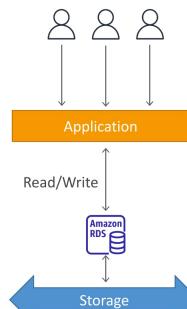


- RDS is a managed service:
 - Automated provisioning, OS patching
 - Continuous backups and restore to specific timestamp (Point in Time Restore)
 - Monitoring dashboards
 - Read replicas for improved read performance
 - Multi AZ setup for DR (Disaster Recovery)
 - Maintenance windows for upgrades
 - Scaling capability (vertical and horizontal)
 - Storage backed by EBS (gp2 or io1)
- BUT you can't SSH into your instances

Don't have access to underlying EC2 instance

RDS – Storage Auto Scaling

- Helps you increase storage on your RDS DB instance dynamically
- When RDS detects you are running out of free database storage, it scales automatically
- Avoid manually scaling your database storage
- You have to set **Maximum Storage Threshold** (maximum limit for DB storage)
- Automatically modify storage if:
 - Free storage is less than 10% of allocated storage
 - Low-storage lasts at least 5 minutes
 - 6 hours have passed since last modification
- Useful for applications with unpredictable workloads
- Supports all RDS database engines



Read Replicas

+ each endpoint acts like a new DNS

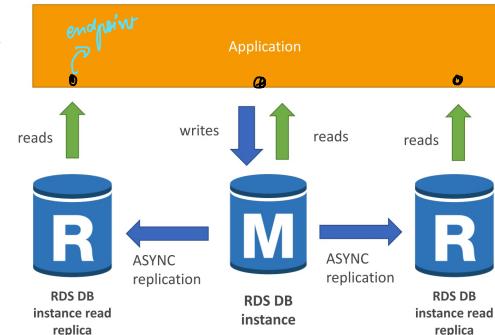
RDS Read Replicas for read scalability

- useful to distribute in read workloads using an asynchronous replication strategy!
- choose endpoint for each replica hence we need to

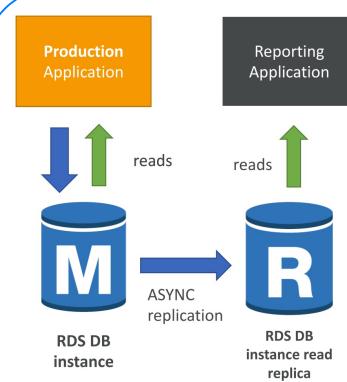
* need not do with multiple AZ!

Use Care!

- Up to 15 Read Replicas
- Within AZ, Cross AZ or Cross Region
- Replication is ASYNC, so reads are eventually consistent
- Replicas can be promoted to their own DB
- Applications must update the connection string to leverage read replicas



- You have a production database that is taking on normal load
- You want to run a reporting application to run some analytics
- You create a Read Replica to run the new workload there
- The production application is unaffected
- Read replicas are used for SELECT (=read) only kind of statements (not INSERT, UPDATE, DELETE)

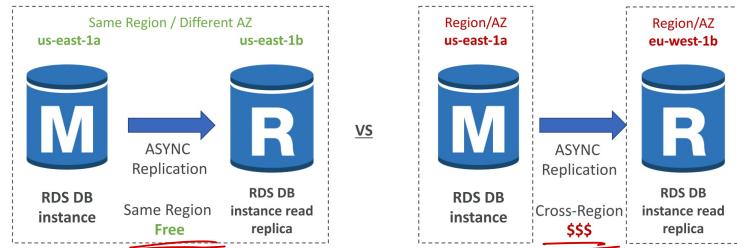


just so that we don't overload the main DB
also
security pov
i.e. only "select" statements!

* generally for managed service with your data transfer may have encryption because in AWS moving data from 1 AZ to other costs!

RDS Read Replicas – Network Cost

- In AWS there's a network cost when data goes from one AZ to another
- For RDS Read Replicas within the same region, you don't pay that fee

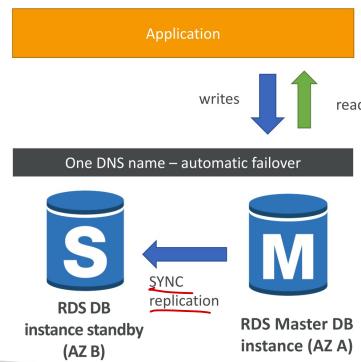


Multi AZ Setup!

→ Read replicas can be used as Multi AZ for disaster recovery but remember that read replicas use ASYNC replication and if the Main DB is being written and suppose a failover occurs, then the older data in read replica will be used.

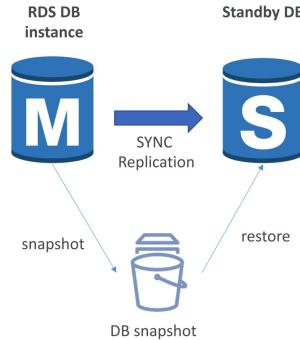
RDS Multi AZ (Disaster Recovery)

- SYNC replication
- One DNS name – automatic app failover to standby
- Increase availability
- Failover in case of loss of AZ, loss of network, instance or storage failure
- No manual intervention in apps
- Not used for scaling
- Note: The Read Replicas be setup as Multi AZ for Disaster Recovery (DR)



RDS – From Single-AZ to Multi-AZ

- Zero downtime operation (no need to stop the DB)
- Just click on "modify" for the database
- The following happens internally:
 - A snapshot is taken
 - A new DB is restored from the snapshot in a new AZ
 - Synchronization is established between the two databases



download for windows!

* Create a database from console

→ use SQL client to connect to the RDS database endpoint so that we can perform query operations on the DB!

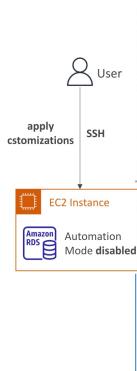
→ to make read replicas create from action.

→ check other actions available!

RDS Custom

→ Allows us to manage the underlying database & O.S which RDS automatically does!

- Managed Oracle and Microsoft SQL Server Database with OS and database customization
 - RDS: Automates setup, operation, and scaling of database in AWS
 - Custom: access to the underlying database and OS so you can
 - Configure settings
 - Install patches
 - Enable native features
 - Access the underlying EC2 Instance using SSH or SSM Session Manager
 - De-activate Automation Mode to perform your customization, better to take a DB snapshot before
 - RDS vs. RDS Custom
 - RDS: entire database and the OS to be managed by AWS
 - RDS Custom: full admin access to the underlying OS and the database



AWS Errors { Aws proprietary }

→ AWS Managed PostgreSQL & MySQL DB
→ also has a newerless implementation Amazon Aurora

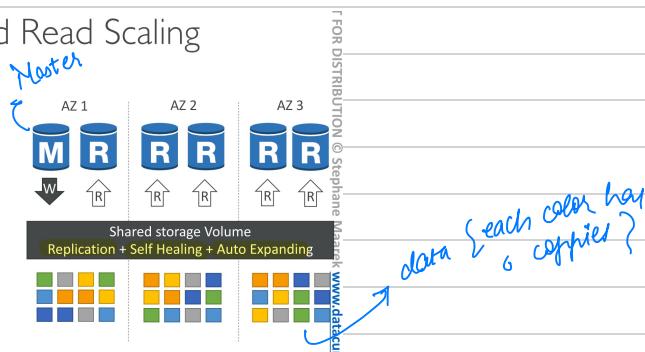
Amazon Aurora



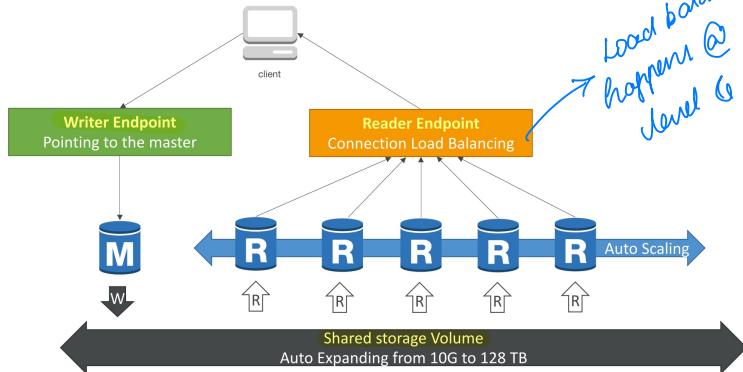
- Aurora is a proprietary technology from AWS (not open sourced)
 - Postgres and MySQL are both supported as Aurora DB (that means your drivers will work as if Aurora was a Postgres or MySQL database)
 - Aurora is “AWS cloud optimized” and claims 5x performance improvement over MySQL on RDS, over 3x the performance of Postgres on RDS
 - Aurora storage automatically grows in increments of 10GB, up to 128 TB.
 - Aurora can have up to 15 replicas and the replication process is faster than MySQL (sub 10 ms replica lag)
 - Failover in Aurora is instantaneous. It’s HA (High Availability) native.
 - Aurora costs more than RDS (20% more) – but is more efficient

Aurora High Availability and Read Scaling

- 6 copies of your data across 3 AZs:
 - 4 copies out of 6 needed for writes
 - 3 copies out of 6 need for reads
 - **Self healing with peer-to-peer replication**
 - Storage is striped across 100s of volumes
 - One Aurora Instance takes writes (master)
 - Automated failover for master in less than 30 seconds
 - Master + up to 15 Aurora Read Replicas serve reads = **16**
 - Support for **Cross Region Replication**



Aurora DB Cluster



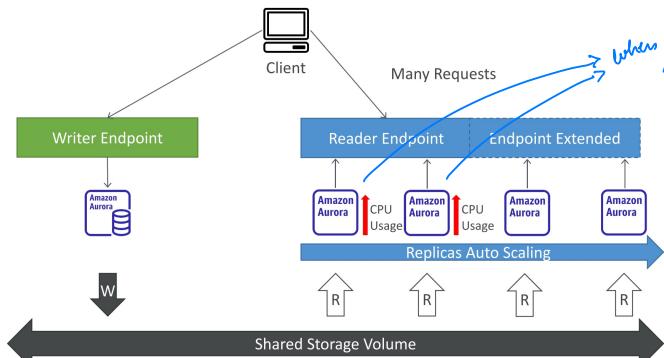
Load balancing
happens @ connection level
not statement level

- * connection level
 - physical hardware
 - client links
- * statement level
 - they execute only once for each single transaction

Features of Aurora

- Automatic fail-over
- Backup and Recovery
- Isolation and security
- Industry compliance
- Push-button scaling
- Automated Patching with Zero Downtime
- Advanced Monitoring
- Routine Maintenance
- **Backtrack:** restore data at any point of time without using backups

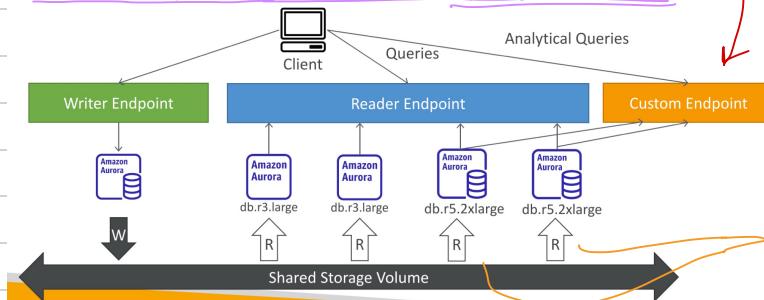
Aurora Replicas - Auto Scaling



when CPU usage for slave nodes increase
new reader endpoint will create replica as an action of the auto scaling feature!

Aurora – Custom Endpoints

- Define a **subset of Aurora Instances** as a **Custom Endpoint**
- Example: Run analytical queries on specific replicas
- The Reader Endpoint is generally not used after defining Custom Endpoints

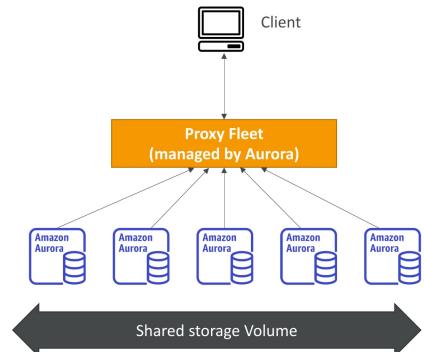


→ think we will be having different custom endpoints for different type of queries!.

larger read replicas

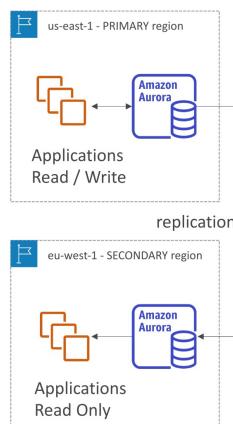
Aurora Serverless

- Automated database instantiation and auto-scaling based on actual usage
- Good for infrequent, intermittent or unpredictable workloads
- **No capacity planning needed**
- Pay per second, can be more cost-effective



Global Aurora

- Aurora Cross Region Read Replicas:
 - Useful for disaster recovery
 - Simple to put in place
- Aurora Global Database (recommended):
 - 1 Primary Region (read / write)
 - Up to 5 secondary (read-only) regions, replication lag is less than 1 second
 - Up to 16 Read Replicas per secondary region
 - Helps for decreasing latency
 - Promoting another region (for disaster recovery) has an RTO of < 1 minute
 - Typical cross-region replication takes less than 1 second



Aurora Machine Learning

- Enables you to add ML-based predictions to your applications via SQL
- Simple, optimized, and secure integration between Aurora and AWS ML services
- Supported services
 - Amazon SageMaker (use with any ML model)
 - Amazon Comprehend (for sentiment analysis)
- You don't need to have ML experience
- Use cases: fraud detection, ads targeting, sentiment analysis, product recommendations



SQL query → what are the recommended products!
take user's data from Aurora & make predictions

RDS Backups

Automated Backups!

- * Daily backup of the DB
 - * earliest backup of logs is 5 minutes in RDS
 - * we can set the automated backup retention period from 1 to 35 days! {so we can disable automated backup}
- gives the ability to restore it to a point in time oldest backup to 5min ago.

Manual DB snapshots

- * user triggered & no retention period problems!

Trick → suppose we need to use the RDS database for just 2 hrs in a month, instead of stopping the DB {which will still cost money because storage is being used} we can take the snapshot of DB & auto the DB, whenever needed restore it!
* taking a snapshot costs much lesser than stopping the entire DB!

Aurora Backups

Automatic Backup

- * 1 - 35 days retention { Cannot be disabled }
- * point in time recovery in that timeframe

Manual Backup

- * same

Restoring RDS or Aurora

- * restoration creates a new DB

Restoring MySQL RDS/Aurora DB from S3

- 1) Take a backup of the on-premises DB cluster
- 2) Store it in an Amazon S3 bucket
- 3) Restore the backup into a new RDS/Aurora cluster running MySQL.

use Percona XtraBackup
to create an Aurora cluster

Aurora Database Cloning

- Create a new Aurora DB Cluster from an existing one
- Faster than snapshot & restore
- Uses copy-on-write protocol
 - ① Initially, the new DB cluster uses the same data volume as the original DB cluster (fast and efficient – no copying is needed)
 - ② When updates are made to the new DB cluster data, then additional storage is allocated and data is copied to be separated
- Very fast & cost-effective
- Useful to create a "staging" database from a "production" database without impacting the production database

→ suppose we need to run some tests on the Aurora DB in its stage the DB away from the production env.

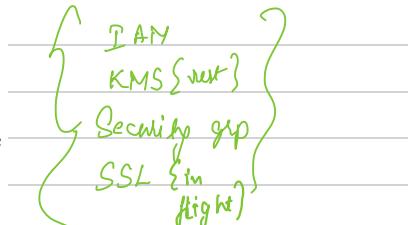


→ its a quick & cost effective feature to copy a database whereas snapshots are read only copy of DB's.

RDS & Aurora Security

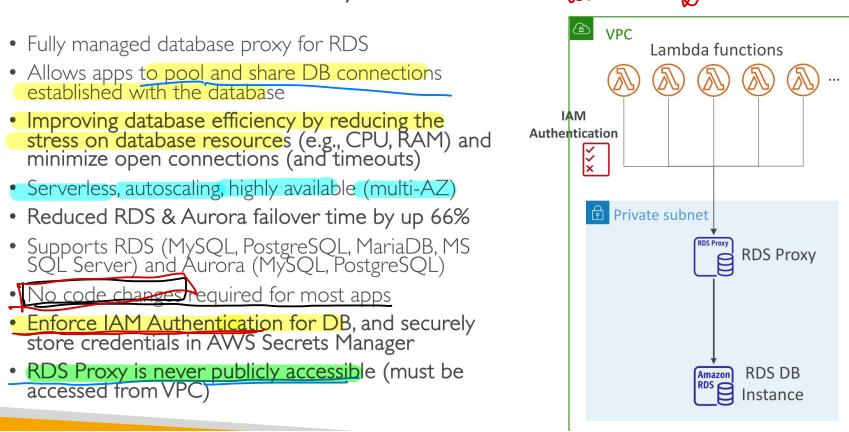
→ PostgreSQL & MySQL support IAM authentication but oracle does not

- At-rest encryption:
 - Database master & replicas encryption using AWS KMS – must be defined as launch time
 - If the master is not encrypted, the read replicas cannot be encrypted
 - To encrypt an un-encrypted database, go through a DB snapshot & restore as encrypted
- In-flight encryption: TLS-ready by default, use the AWSTLS root certificates client-side
- IAM Authentication: IAM roles to connect to your database (instead of username/pw)
- Security Groups: Control Network access to your RDS / Aurora DB
- No SSH available except on RDS Custom
- Audit Logs can be enabled and sent to CloudWatch Logs for longer retention



Amazon RDS Proxy

- Fully managed database proxy for RDS
- Allows apps to pool and share DB connections established with the database
- Improving database efficiency by reducing the stress on database resources (e.g., CPU, RAM) and minimize open connections (and timeouts)
- Serverless, autoscaling, highly available (multi-AZ)
- Reduced RDS & Aurora failover time by up 66%
- Supports RDS (MySQL, PostgreSQL, MariaDB, MS SQL Server) and Aurora (MySQL, PostgreSQL)
- ~~No code changes required for most apps~~
- Enforce IAM Authentication for DB, and securely store credentials in AWS Secrets Manager
- RDS Proxy is never publicly accessible (must be accessed from VPC)



→ Reduces failures
How by 66%.

→ A funcⁿ can multiply if they may have open connections or timeouts to the DB so connecting it to a proxy will prevent the unloading of DB instance. This may overload the proxy [but it is meant to be]

Amazon ElastiCache Overview



→ Results of queries are cached!

→ Using the ever version makes it stateless

- The same way RDS is to get managed Relational Databases...
- ElastiCache is to get managed Redis or Memcached
- Caches are in-memory databases with really high performance, low latency
- Helps reduce load off of databases for read intensive workloads
- Helps make your application stateless
- AWS takes care of OS maintenance / patching, optimizations, setup, configuration, monitoring, failure recovery and backups
- Using ElastiCache involves heavy application code changes

stateless application

What each request as independent transaction that can be handled by any available servers.

→ It does not use client data in one session for up in next version.

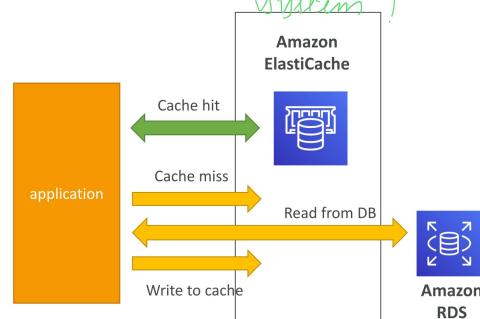
ElastiCache

Solution Architecture - DB Cache

- Applications queries ElastiCache, if not available, get from RDS and store in ElastiCache.
- Helps relieve load in RDS
- Cache must have an invalidation strategy to make sure only the most current data is used in there.

using RDS proxy which does not require code changes

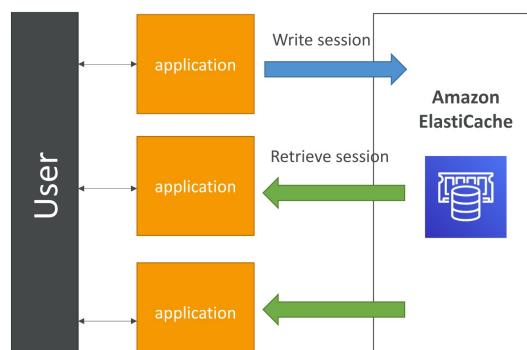
just how the cache works in every other system!



ElastiCache

Solution Architecture – User Session Store

- User logs into any of the application
- The application writes the session data into ElastiCache
- The user hits another instance of our application
- The instance retrieves the data and the user is already logged in



→ for session restoration!

→ Stateless because the data is not stored in the EC2 servers.

ElastiCache – Redis vs Memcached

REDIS

- Multi AZ with Auto-Failover
- Read Replicas to scale reads and have high availability
- Data Durability using AOF persistence
- Backup and restore features
- Supports Sets and Sorted Sets



MEMCACHED

- Multi-node for partitioning of data (sharding)
- No high availability (replication)
- Non persistent
- No backup and restore
- Multi-threaded architecture



AOF persistence {Append only File}

→ logs every write operation
received by the server to an AOF
allows point-in-time recovery!

→ AOF provides better durability
than RDB {Redo Data Base file}
persistence, which takes the
snapshot of the dataset at
specific intervals!

→ AOF allows entire dataset
recovey in the time of crash
it also doesn't block current
write operations.

ElastiCache – Cache Security

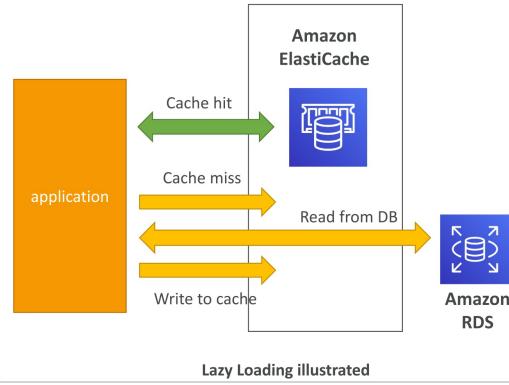
- ElastiCache supports **IAM Authentication for Redis only**
- IAM policies on ElastiCache are only used for AWS API-level security
- Redis AUTH → Redis already has this!
 - You can set a "password/token" when you create a Redis cluster
 - This is an extra level of security for your cache (on top of security groups)
 - Support SSL in flight encryption
- Memcached
 - Supports SASL-based authentication (advanced)

Simple authentication on a memory layer.



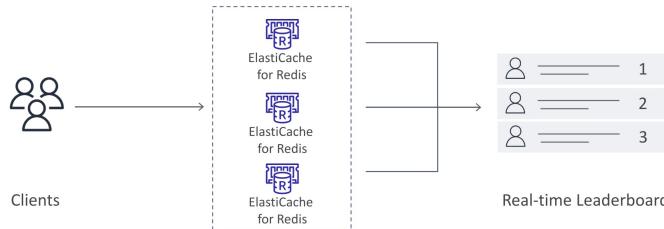
Patterns for ElastiCache

- ① • **Lazy Loading:** all the read data is cached, data can become stale in cache
- ② • **Write Through:** Adds or update data in the cache when written to a DB (no stale data)
- ③ • **Session Store:** store temporary session data in a cache (using TTL features)
 - Quote: There are only two hard things in Computer Science: cache invalidation and naming things



ElastiCache – Redis Use Case

- **Gaming Leaderboards** are computationally complex
- Redis Sorted sets guarantee both uniqueness and element ordering
- Each time a new element added, it's ranked in real time, then added in correct order



→ with a Redis cluster we can create a dual-time leaderboard which use Redis sorted sets & also ordered each time new current is added

→ We don't need entire

Important ports

W

RDS DB ports

FTP - 21

HTTP - 80

HTTPS - 443

SSH - 22

FSTP - 22

PostgreSQL - 5432

MySQL - 3306

Oracle RDS - 1521

MSSQL Server - 1433

MariaDB - 3306

Aurora - 5432 [if postgres]

or
3306 [if MySQL]

Databases in AWS

Database Types

online transaction processing



online analytic processing!

- RDBMS (= SQL / OLTP): RDS, Aurora – great for joins
 - NoSQL database – no joins, no SQL : DynamoDB (~JSON), ElastiCache (key / value pairs), Neptune (graphs), DocumentDB (for MongoDB), Keyspaces (for Apache Cassandra)
 - Object Store: S3 (for big objects) / Glacier (for backups / archives)
 - Data Warehouse (= SQL Analytics / BI): Redshift (OLAP), Athena, EMR
 - Search: OpenSearch (JSON) – free text, unstructured searches
 - Graphs: Amazon Neptune – displays relationships between data
 - Ledger: Amazon Quantum Ledger Database
 - Time series: Amazon Timestream
- Note: some databases are being discussed in the Data & Analytics section

→ data durability :-
data guarding data
from loss or corruption
in the event of
failures or outages

Choosing the Right Database

- We have a lot of managed databases on AWS to choose from
- Questions to choose the right database based on your architecture:
 - Read-heavy, write-heavy, or balanced workload? Throughput needs? Will it change, does it need to scale or fluctuate during the day?
 - How much data to store and for how long? Will it grow? Average object size? How are they accessed?
 - Data durability? Source of truth for the data ?
 - Latency requirements? Concurrent users?
 - Data model? How will you query the data? Joins? Structured? Semi-Structured?
 - Strong schema? More flexibility? Reporting? Search? RDBMS / NoSQL?
 - License costs? Switch to Cloud Native DB such as Aurora?

Amazon RDS – Summary



- Managed PostgreSQL / MySQL / Oracle / SQL Server / DB2 / MariaDB / Custom
 - Provisioned RDS Instance Size and EBS Volume Type & Size
 - Auto-scaling capability for Storage
 - Support for Read Replicas and Multi AZ
 - Security through IAM, Security Groups, KMS ,SSL in transit
 - Automated Backup with Point in time restore feature (up to 35 days)
 - Manual DB Snapshot for longer-term recovery
 - Managed and Scheduled maintenance (with downtime)
 - Support for IAM Authentication, integration with Secrets Manager
 - RDS Custom for access to and customize the underlying instance (Oracle & SQL Server)
- Use case: Store relational datasets (RDBMS / OLTP), perform SQL queries, transactions

Amazon Aurora – Summary



- Compatible API for PostgreSQL / MySQL, separation of storage and compute
 - Storage: data is stored in 6 replicas, across 3 AZ – highly available, self-healing, auto-scaling *default?*
 - Compute: Cluster of DB Instance across multiple AZ, auto-scaling of Read Replicas
 - Cluster: Custom endpoints for writer and reader DB instances
 - Same security / monitoring / maintenance features as RDS
 - Know the backup & restore options for Aurora
 - Aurora Serverless – for unpredictable / intermittent workloads, no capacity planning
 - Aurora Global: up to 16 DB Read Instances in each region, < 1 second storage replication
 - Aurora Machine Learning: perform ML using SageMaker & Comprehend on Aurora
 - Aurora Database Cloning: new cluster from existing one, faster than restoring a snapshot
 - Use case: same as RDS, but with less maintenance / more flexibility / more performance / more features
- continuous backup to S3*

Amazon ElastiCache – Summary



- Managed Redis / Memcached (similar offering as RDS, but for caches)
 - In-memory data store, sub-millisecond latency
 - Select an ElastiCache instance type (e.g. cache.m6g.large)
 - Support for Clustering (Redis) and Multi AZ, Read Replicas (sharding)
 - Security through IAM, Security Groups, KMS, Redis Auth
 - Backup / Snapshot / Point in time restore feature
 - Managed and Scheduled maintenance
 - Requires some application code changes to be leveraged
- Use Case: KeyValue store, frequent reads, less writes, cache results for DB queries, store session data for websites, cannot use SQL.

Fully managed, high available, NoSQL database with replication across 3 AZ. Key-value pair database. Flagship AWS product which scales to massive workloads, distributed serverless database. (with rds or elastic cache we need to provide a instance type) still servers at backend. Millions of request per seconds, trillions of rows 100s of TB of storage

*by dynamoDB
dynamodb
in memory
cache for dynamo DB*

Amazon DynamoDB – Summary

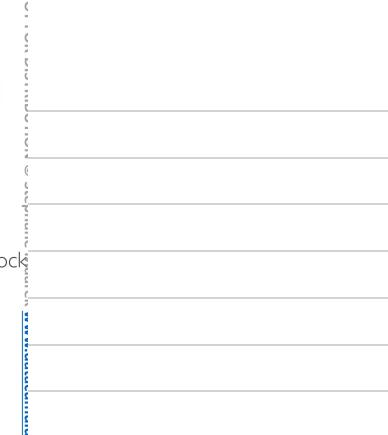


- AWS proprietary technology, managed serverless NoSQL database, millisecond latency
 - Capacity modes: provisioned capacity with optional auto-scaling or on-demand capacity
 - Can replace ElastiCache as a key/value store (storing session data for example, using TTL feature)
 - Highly Available, Multi AZ by default, Read and Writes are decoupled, transaction capability
 - DAX cluster for read cache, microsecond read latency
 - Security, authentication and authorization is done through IAM *to provide a secure interface to the dynamoDB cluster*
 - Event Processing: DynamoDB Streams to integrate with AWS Lambda, or Kinesis Data Streams
 - Global Table feature: active-active setup *two multi-region applications*
 - Automated backups up to 35 days with PITR (restore to new table), or on-demand backups
 - Export to S3 without using RCU within the PITR window, import from S3 without using WCU
 - Great to rapidly evolve schemas *Read capacity unit!*
- Use Case: Serverless applications development (small documents 100s KB), distributed serverless cache *with capacity unit!*

Amazon S3 – Summary



- S3 is a... key / value store for objects
- Great for bigger objects, not so great for many small objects
- Serverless, scales infinitely, max object size is 5 TB, versioning capability
- Tiers: S3 Standard, S3 Infrequent Access, S3 Intelligent, S3 Glacier + lifecycle policy
- Features: Versioning, Encryption, Replication, MFA-Delete, Access Logs...
- Security: IAM, Bucket Policies, ACL, Access Points, Object Lambda, CORS, Object/Vault Lock
- Encryption: SSE-S3, SSE-KMS, SSE-C, client-side, TLS in transit, default encryption
- Batch operations on objects using S3 Batch, listing files using S3 Inventory
- Performance: Multi-part upload, S3 Transfer Acceleration, S3 Select
- Automation: S3 Event Notifications (SNS, SQS, Lambda, EventBridge)
- Use Cases: static files, key value store for big files, website hosting



→ for document format
of NoSQL database/
store in the form of
JSON document

DocumentDB NoSQL DB



- Aurora is an "AWS-implementation" of PostgreSQL / MySQL ...
- DocumentDB is the same for MongoDB (which is a NoSQL database)
- MongoDB is used to store, query, and index JSON data
- Similar "deployment concepts" as Aurora
- Fully Managed, highly available with replication across 3 AZ
- DocumentDB storage automatically grows in increments of 10GB
- Automatically scales to workloads with millions of requests per seconds

JSON

JSON

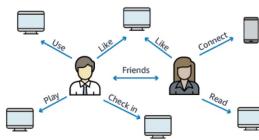
```
1  [
2    {
3      "year": 2013,
4      "title": "Turn It Down, Or Else!",
5      "info": {
6        "directors": ["Alice Smith", "Bob Jones"],
7        "release_date": "2013-01-18T00:00:00Z",
8        "rating": 6.2,
9        "genres": ["Comedy", "Drama"],
10       "image_url": "http://ia.media-imdb.com/images/N/09ERWf
11       "plot": "A rock band plays their music at high volumes
12       "actors": ["David Matthewman", "Jonathan G. Neff"]
13     }
14   },
15   {
16     "year": 2015,
17     "title": "The Big New Movie",
18     "info": {
19       "plot": "Nothing happens at all.",
20       "rating": 0
21     }
22   }
23 ]
```

Amazon Neptune { graph database }



→ AWS managed
graph database
a NoSQL DB

- Fully managed graph database
- A popular graph dataset would be a social network
 - Users have friends
 - Posts have comments
 - Comments have likes from users
 - Users share and like posts...
- Highly available across 3 AZs with up to 15 read replicas
- Build and run applications working with highly connected datasets – optimized for these complex and hard queries
- Can store up to billions of relations and query the graph with milliseconds latency
- Highly available with replications across multiple AZs
- Great for knowledge graphs (Wikipedia), fraud detection, recommendation engines, social networking

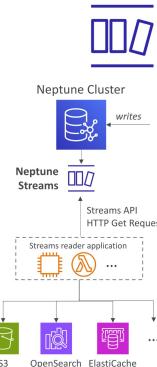


→ Any change to
the DB is recorded
& can be reviewed
& enacted upon
immediately, like
sending notification
if someone likes a
post, etc.

Amazon Neptune – Streams

- Real-time ordered sequence of every change to your graph data
- Changes are available immediately after writing
- No duplicates, strict order
- Streams data is accessible in an HTTP REST API

Application can access
the stream data
using REST API



Amazon Keyspaces (for Apache Cassandra)



Managed
Apache Cassandra

- Apache Cassandra is an open-source NoSQL distributed database
- A managed Apache Cassandra-compatible database service
- Serverless, Scalable, highly available, fully managed by AWS
- Automatically scale tables up/down based on the application's traffic
- Tables are replicated 3 times across multiple AZ
- Using the Cassandra Query Language (CQL)
- Single-digit millisecond latency at any scale, 1000s of requests per second
- Capacity: On-demand mode or provisioned mode with auto-scaling
- Encryption, backup, Point-In-Time Recovery (PITR) up to 35 days
- Use cases: store IoT devices info, time-series data, ...

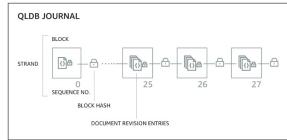
Similar to
DynamoDB

Cassandra is a
NoSQL
Open Source
distributed
column store
Cayenne of Key-Value
What's a database?

Amazon QLDB



- QLDB stands for "Quantum Ledger Database"
- A ledger is a book recording financial transactions
- Fully Managed, Serverless, High available, Replication across 3 AZ
- Used to review history of all the changes made to your application data over time
- Immutable system: no entry can be removed or modified, cryptographically verifiable



- 2-3x better performance than common ledger blockchain framework
- Difference with Amazon Managed Blockchain: no decentralization component, in accordance with financial regulation rules

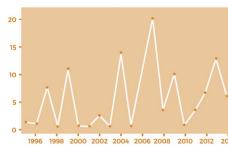
→ Behind the scenes there is a journal which has a seq. of mod. so any time an modification is made there is a cryptographic hash which is computed

That guarantees that nothing has been deleted or changed.

Amazon Timestream {time series DB}



specificity for time series data because using a relational DB for it is very costly.



→ from the readings & do just analytics.

- Fully managed, fast, scalable, serverless time series database
- Automatically scales up/down to adjust capacity
- Store and analyze trillions of events per day
- 1000s times faster & 1/10th the cost of relational databases
- Scheduled queries, multi-measure records, SQL compatibility
- Data storage tiering: recent data kept in memory and historical data kept in a cost-optimized storage
- Built-in time series analytics functions (helps you identify patterns in your data in near real-time)
- Encryption in transit and at rest

- Use cases: IoT apps, operational applications, real-time analytics, ...

Amazon Timestream – Architecture

